

Limitations of the GradCafe Dataset

The GradCafe dataset has important limitations because the information is voluntarily submitted by users rather than collected from a random or complete sample of graduate school applicants. Applicants who choose to report their results may differ from applicants who do not, so acceptance rates, GPA averages, and other statistics calculated from this dataset should not be interpreted as representative of all graduate applicants. In addition, many fields are self-reported, which introduces the possibility of entry errors, inconsistent formatting, or inaccurate values. For example, my analysis produced an average GRE Quantitative score of 259.81, which indicates that some reported values are outside the expected GRE scale. This demonstrates why summary statistics from the dataset should be interpreted cautiously rather than treated as official admissions statistics.

A second limitation is that the dataset contains missing and inconsistently formatted information. School names, program names, degree types, citizenship status, GPA, and standardized test scores may be absent or entered differently by different users. This can affect filtering, grouping, and comparisons across programs or universities. The effect is visible in my analysis of accepted Fall 2026 PhD Computer Science applicants at Georgetown, MIT, Stanford, and Carnegie Mellon: using the original fields returned 28 records, while using the LLM-generated standardized program and university fields returned 24. The difference shows that data cleaning and LLM-based standardization can change query results and may introduce classification errors of their own. Therefore, the results are useful for exploring patterns in the GradCafe dataset, but they should not be interpreted as definitive measures of university admissions outcomes.